

Lforge: Extending Forge with an Interactive Theorem Prover

Embedding the Forge specification language via a machine translation
as a language-level feature of the Lean theorem prover.

Jiahua Chen

Advisor: Robert Lewis

Reader: Tim Nelson

A Thesis submitted in partial fulfillment of the requirements for Honors
in the Departments of Mathematics and Computer Science at Brown University



Department of Computer Science
Brown University
Providence, Rhode Island
April, 2024

Abstract

While formal methods are being increasingly applied in industry and are invaluable in allowing users to specify, model, and verify complex systems, the multitude of available tools offer little to no interoperability amongst each other. There are no common formal methods models nor specification languages, and each tool offers precisely what its specific logical framework allows. This creates a dilemma in the field that while users understand the strengths and capabilities of formal methods tools, they are often not able to select a tool appropriate for the task at hand. We present LFORGE, a tool that implements the Forge specification language via a translation process as a language-level feature of Lean. LFORGE offers a ‘best-of-both-worlds’ approach—allowing users to no longer be constrained by the finite yet automatic resolving capabilities of Forge nor the tedious proof process of Lean. LFORGE serves as an example of interfacing between two drastically different tools, thus allowing users to harness the resolving capabilities of two frameworks: Forge, an offshoot of the Alloy specification language based on a relational logic solver that can automatically prove facts about finite models; and Lean, an interactive theorem prover that allows users to prove generalized statements. In doing so, LFORGE serves as a model of what specification portability might look like across different classes of formal methods tools, as well as a model of user experience in such a translation. LFORGE is furthermore one of the first experimentations of Lean 4’s rich metaprogramming capabilities. This allows LFORGE to be a full-fledged Forge DSL within Lean with a focus on usability-first despite operating within the constraints imposed by Forge and Lean’s respective formal frameworks.

Contents

Notation	iii
1. Introduction	1
2. Background	3
2.1. Related Work	3
2.2. <i>Lean</i> and other proof assistants	3
2.3. <i>Forge</i> , <i>Alloy</i> , and other relational specification languages	3
3. Motivation	4
3.1. A Toy Example	4
4. Design	7
4.1. Design Summary	7
4.2. Syntax, Parsing, and the Forge AST	7
4.3. Elaboration	9
4.4. The Forge Model in Lean	9
5. Implementation Details and Challenges	10
5.1. Finiteness of Forge Sigs	10
5.2. Sig Inheritance and Quantifiers	11
5.3. “Everything is a Set”	13
5.4. Type Coercions	16
5.5. Relational Joins, Cross, Inclusion, Equality	16
5.6. Integers	16
6. Results and Examples	20
6.1. Forge as a Lean DSL	20
6.2. A Toy Example, <i>Continued</i>	25
6.3. A Mutual-Exclusion Protocol	25
7. Discussion	26
7.1. Contributions	26
7.2. Reflections	26
7.3. Future Work	26
Bibliography	29
Appendices	30

Notation

Throughout this paper, you might encounter some visual or symbolic notation that is specific to this project or specific to the semantics of Forge and/or Lean. These are noted here.

“ \bowtie ” denotes the relational join operator. If $x : A \rightarrow B$ is a relation and $y : B \rightarrow C$ is a relation, then $x \bowtie y$ produces the relation $A \rightarrow C$ merged on common values in the rightmost (B) column of x and the leftmost (B) column of y . x and y can be of arbitrary arity, so long as their leftmost and rightmost columns respectively match.

Code snippets and listings have been included in this paper to serve as examples, motivation, or to provide implementation details. Where they are included, the color of the code block denotes the source language and context.

fix according
to feedback

```
1 -- This is the code block for the Lean implementation of our translation
2 def forgeEnsureHasType (expectedType? : Option Expr) (e : Expr)
3   (errorMsgHeader? : Option String := "Forge Type Error")
4   (f? : Option Expr := none) : TermElabM Expr := do
5   let some expectedType := expectedType? | return e
6   if (← isDefEq (← inferType e) expectedType) then
7     return e
8   else
9     mkCoe expectedType e f? errorMsgHeader?
```

is an example of a Lean implementation code block. This denotes code from the implementation of the translation from Forge to Lean. This encompasses the parsing and elaboration of Forge syntax within Lean, and is most often the metaprogramming implementation of Forge in Lean.

```
1 -- This is the code block for a snippet of a model specification in Forge
2 sig Node {
3   neighbors : set Node
4 }
5 pred connected[a : Node, b : Node] {
6   b in a.neighbors
7 }
```

is an example of a Forge code block. This denotes examples of a model (or a snippet of a model) in Forge.

```
1 -- This is the code block for the translated Lean equivalent of a Forge snippet
2 opaque Node : Type
3 opaque neighbors : Node → Node → Prop
4
5 def connected (a : Node) (b : Node) : Prop :=
6   neighbors a b
```

is an example of a Lean translation code block. This denotes examples of the translated version of a Forge model or snippet. This is oftentimes the translated Lean code that is emitted out of our program.

1. Introduction

Formal methods are increasingly being applied in industry. Domain-specific formal methods tools empower researchers to specify, model, analyze, and verify complex software and hardware systems that otherwise prove unfeasible to fully examine by hand [1, 18]. These applications prove invaluable when the functionality of an existing system needs to be verified to be correct, or when new systems need to be synthesized based on a set of logical constraints and specifications [33].

Yet, there is a multitude of tools and flavors that exist in the realm of formal methods: type-checked programming languages [12, 5], property-based testing frameworks [11, 20], modeling and specification languages [15, 14, 28], SMT solvers [7], proof assistants [24], etc. Each of these tools offers a tailored set of features and is based upon specific yet different logical framework. An SMT solver based on the boolean satisfiability problem is different from a proof assistant which relies on dependent type theory. Due to operating under these diverse and different frameworks, few if any offer any form of interoperability—there is no common form of formal methods modeling or specification language. Each tool offers precisely what its framework allows.

add citation?

As a result, there arise limitations as to what *can* and *cannot* be modeled by specific techniques. In a survey of applications of formal methods within industry, the majority of respondents repeatedly identified that formal methods are useful in projects but similarly agreed that they felt tools were incapable or often ill-suited to the particular task at hand [33]. This problem of selecting tools is so prevalent that there have been surveys and methodologies devised for this task [17, 6].

This paper introduces LFORGE¹, a tool that implements the Forge specification language [28] (a pedagogical offshoot of Alloy [15], which we use due to its gradually featured nature as well as simpler syntax) via a translation process as a language-level feature of Lean, an interactive proof assistant [24].

In doing such, LFORGE aims to serve as an example of interfacing between two drastically different tools in the larger realm of formal methods. The goal of this is to reduce the number of ‘make-or-break’ choices that researchers face within the field and allow users to harness the resolving capabilities of multiple formal methods models, picking and choosing the features from multiple feature sets that are important to them.

Forge and Lean work in fundamentally different ways. Forge, which is based on the Alloy relational model solver, uses the language of relational logic to specify and ‘solve’ systems. A system is specified as a collection of *signatures*, and model specifications are provided as a set of logical constraints on relations between signatures. Forge will then apply a SAT solver to the set of constraints to generate an *instance* of the specified model (with some finite instances of each signature) [14, 28]. This makes Forge suited for generating finite examples or counterexamples to provided specifications. Lean, on the other hand, is an interactive theorem prover that is based on dependent type theory² [3]. Systems are implemented within the functional programming language, and

¹Unfortunately the superior option amongst contenders FLEAN, FEAN, and LORGE.

²Specifically, the *calculus of (inductive) constructions*. This is the logical system first implemented in Coq [4].

theorems containing logical statements about said systems can be stated and proven. Lean verifies that said proof is correct—and that the system has claimed properties.

While Forge can automatically reason (via solving for satisfying instances) about finite instances and can solve for model existence, Lean allows the user to make general claims about a system, at the cost of requiring manual proving. Lean can assist in generalizing statements made in Forge, while Forge can easily disprove incorrect Lean statements via counterexample³. By allowing users to input Forge directly into a Lean source program, users can on one hand harness the automated reasoning tools that Forge provides, yet circumvent any bound limitations of Forge by proving theorems directly in Lean. LFORGE recognizes the benefit of being able to interoperate between these two models for verifying systems can prove useful in checking real-world models, especially where a human translation between the two tools can be tedious and prone to errors.

The implication of this work and further hope is that the model specification syntax of Forge/Alloy can become a universal and portable specification suitable for multiple tools based on multiple frameworks alike. While we do make compromises as to what is and isn't able to be brought over from Forge to Lean, LFORGE is explicit and clear about those assumptions and what is left for the user to specify or supplement. The larger objective is for LFORGE to serve as a model of what specification portability could look like across classes of formal methods tools, and what the user experience might be as specifications are being translated and utilized.

Furthermore, LFORGE serves as one of the first experimentations with Lean 4's complex and rich metaprogramming capabilities [29], implementing Forge as a full-fledged DSL in Lean 4. At the same time, we harness Lean 4's out-of-the-box language server⁴, interactive capabilities [25], and type unification system to suit Forge. The end product is a language experience that is definitely on par, if not more feature-rich than native Forge support in any IDEs from a purely user-experience perspective.

The hope is that LFORGE becomes a DSL that focuses on usability-first despite operating within the translation constraints imposed by formal frameworks. We hope for it to become a tool (or, at the very least, serve as a blueprint for a tool) that is an essential utility in any researcher's formal methods toolbox.

³Lean will not indicate whether a statement is true or false.

⁴What powers Lean 4's autocomplete, interactive theorem prover, tooltip-on-hover, etc. in VS Code.

2. Background

2.1. Related Work

The following is a summary of some of the sources and related work I've found: I still need to write this section.

write this
section

- [32] discusses Lean 4's capabilities of extensible syntax.
- [19] describes interoperability between Lean and Mathematica.
- [13] implements the small scale reflection tactic language as a Lean DSL.
- [28] discusses sublanguages of Forge, and that these are viable. We can reasonably restrict language features.
- [22] makes Alloy executable in Java.
- [23] embeds Alloy in Ruby
- [21] not to toot this Milicevic's horn—describes programming environment where imperative and declarative styles are interweaved.
- [9] discuss environments *built around* Forge, for both data processing and verification at simultaneously. *solver-aided programming language*
- [16] specification as data, embed B specification language in Clojure.

2.2. *Lean* and other proof assistants

- [25] discusses Lean user interface
- [24] discusses Lean's extensibility.

2.3. *Forge*, *Alloy*, and other relational specification languages

- Discuss Alloy and Forge, how they work, etc.
- [31] kodkod

3. Motivation

complete

3.1. A Toy Example

We envision the workflow of someone using LFORGE to be as follows:

1. The user specifies a model in interest, or they already have a completed model in Forge that they would like to formalize. Within Forge, the user writes relevant predicates and uses Forge’s SAT backend to quickly check if they are satisfiable (and what satisfying instances look like), or if they are not satisfiable (that is, the negation is a theorem). Users might isolate specific predicates that they would like to prove in detail.
2. In a new Lean file, the user imports the LFORGE module and pastes their Forge specification in, verbatim. If they started the process working in the LFORGE subset of Forge, no changes need to be made. If they are working with an existing Forge file, our tool will prompt the user to make any modifications necessary to keep it compliant with the subset of syntax, including potential type annotations (see [section 5.4](#)). After this, all sigs, fields, predicates are available in Lean.
3. Optionally, a user might want to make *additional* claims or write predicates using Lean’s syntax. They have the option to do so here (see Mixed Execution, [section 6.1](#)).
4. Finally, the user will identify important predicates within their specification that they want to prove generally. They do so using the interactive theorem prover in Lean.

We start by providing a small example that might represent a Forge specification and the desired generated code in Lean as motivation behind this workflow. Bertrand Russell, in illustrating Russell’s paradox on sets, poses the following paradox: “Let the barber be ‘one who shaves all those, and those only, who do not shave themselves.’ The question is, does the barber shave himself?” [30, p. 101].

We might want to construct a formal model for this village for which the barber shaves all those who do not shave themselves, and use Forge to prototype and quickly check properties regarding this model:

```
sig Person {
  shaver: one Person
}

pred shavesThemselves[p: Person] {
  p = p.shaver
}

pred existsBarber {
  some barber : Person | all p : Person | {
    not shavesThemselves[p] <=> p.shaver = barber
  }
}
```


Attempting to run this model for `existsBarber` will yield an `Unsatisfiable`. After all, if the barber doesn't shave themselves, yet they shave all those who don't, they they ought to shave themselves. To prove this statement (and not merely rely on the fact that Forge was not able to find a satisfiable instance within its bounds), we need to transition to a theorem prover.

LFORGE aids in porting the entire specification into Lean, producing a set of equivalent Lean definitions:

```
opaque Person : Type
opaque shaver : Person → Person

def shavesThemselves : Person → Prop :=
  fun p ↦ shaver p p

def existsBarber : Prop :=
  ∃ (barber : Person), ∀ (p : Person), ¬shavesThemselves p ↔ shaver p barber
```

At this point, we might want to continue to provide a mathematical proof, as we observed in Forge, that there cannot possibly be a barber in this town. That is,

```
theorem no_barber : ¬ existsBarber := by ...
```

The conclusion of this example, including the Lean proof of our property, is provided in [section 6.2](#). Note that this was not possible working solely in Forge, which only checks for the existence of examples or counterexamples within the bounds that it knows.

While this was a simple example, it serves as a demonstration of what is possible under this dual framework. One might provide a specification of a protocol and prove that no vulnerabilities exist (or that it has all the desired properties). Lean would serve invaluable in proving the correctness of properties about any real-life system described and prototyped in the Forge specification language. All this stays true to the goal of providing better formal methods tools that are more universally applicable and practical.

Furthermore, the pedagogical implications of such a program is also worth noting. Forge, designed to be a pedagogical language, seeks to teach formal methods gradually and introduce students to formal methods tools enabling them to work better in the real world and industry [28]. Students who are interested in formal methods often take Logic for Systems in the Computer Science department at Brown, the course which Forge was developed for and taught in. Students continue on to take Formal Proof and Verification, a course that teaches the use of proof assistants via Lean. Both courses have open-ended research-style final projects that encourages students to explore and formalize topics that they are interested in. We believe LFORGE also provides an invaluable opportunity for students who have background with both formal methods tools (or who might merely want to explore beyond) to bridge their knowledge between automated reasoning and formalization via proof. While Forge instructs us that a specification or predicate is true, Lean reveals *why* the specification is true. Oftentimes, this makes it difficult to debug incorrect specifications without additional tools like visualizations [28]. By attempting to construct proofs of model properties, stu-

dents are compelled to think about their modeling choices and justify each statement, seamlessly translating between the statements they are making and proofs of their correctness [2].

4. Design

4.1. Design Summary

Lean 4 is a good target for our translation as well as a suitable language to implement our translation in because of the fact that Lean 4 is mostly implemented in itself. As users, we are able to utilize and emit the same data structures used in Lean’s own implementation to extend the functionality of Lean [24]. These metaprogramming capabilities of Lean make our work implementing a Forge module in Lean much easier.

The Lean 4 compilation process is structured as in [fig. 1](#).




Figure 1. A diagram from [29] summarizing the Lean 4 compilation process.

Specifically, the parsing and elaboration steps are designed to be highly customizable and are provided as a ‘first-class’ feature of Lean 4. We approach the problem of translating Forge into Lean as a task of adding new language features to Lean itself. We define new Lean syntax objects that correspond to a concrete syntax tree of Forge and implement a parser for Forge (see [section 4.2](#)), and then implement a custom elaboration function for our Forge syntax to translate it into native Lean expressions (see [section 4.3](#)).

As a result, there is as little additional overhead as possible when translating a Forge specification in Lean. After users have imported our module, all Forge expressions *are* valid Lean expressions and the two languages can be used interchangeably⁵.

complete?

4.2. Syntax, Parsing, and the Forge AST

In the case of parsing, by defining the Forge grammar in the same specification format that Lean defines its syntax in, we can rely on Lean’s parser to parse Forge source code for us.

The benefits of this are two-sided:

1. We are provided `Syntax`-typed Lean objects at the end of this process, the same type that parsing a Lean program would produce. This enables us to treat our Forge implementation as implementation of additional Lean language features, and we can also harness Lean metaprogramming libraries along the way. This is to say, we are implementing Forge in Lean the *same way* Lean is implemented in Lean, which greatly reduces our burden for additional implementation overhead.

⁵We are very fortunate that there are no conflicts between Forge and Lean syntax that would hinder this.

2. By defining Forge ‘blocks’ as a Lean `command`—which is the top-level syntax category⁶—users of the tool can insert raw Forge, without any annotations that this is an “extension language”, into Lean. With the addition of an import statement, every Forge program *is* a valid Lean program.

Lean allows us to create *syntax categories* for each term in our grammar. At the top-level, we have defined `f_sig` (Sigs), `f_pred` (Predicates), `f_fun` (Functions). Terms are either `f_fm1a` for formulas (evaluate to True or False) or `f_expr` for expressions (evaluate to a set, relation, int).

variable?

For example, the grammar⁷ of Forge arguments and predicates is:

$$\begin{aligned}\langle arg \rangle & ::= \langle ident \rangle, + \text{ ‘:’ } \langle expr \rangle \\ \langle args \rangle & ::= \langle arg \rangle, * \\ \langle pred \rangle & ::= \text{ ‘pred’ } \langle ident \rangle [\text{ ‘[’ } \langle args \rangle \text{ ‘]’ }] \text{ ‘{’ } \langle fmla \rangle * \text{ ‘} \text{’ }\end{aligned}$$

Which we can translate into a corresponding syntax definition in Lean:

```
1 declare_syntax_cat f_arg
2 syntax ident, + ":" f_expr : f_arg
3
4 declare_syntax_cat f_args
5 syntax f_arg, * : f_args
6
7 declare_syntax_cat f_pred
8 syntax "pred" ident ("[" f_args "]" )? "{" f_fm1a* "}" : f_pred
```

Following this blueprint, we are able to translate the entirety of the grammar of Forge⁸ into Lean syntax definitions.

What remains to be done is to convert syntax, almost one-to-one, into an AST for Forge, and then *elaborate* (see [section 4.3](#)) our AST into Lean expressions and declarations.

Reference a file/directory in code.

```
1 structure Predicate where
2   name : Symbol
3   name_tok : Syntax
4   args : List (Symbol × Expression) -- (name, type) pairs
5   body : Formula -- with args bound
6   deriving Repr, Inhabited
```

The associated structure definitions of the Forge AST is a deep embedding of Forge into Lean. Using this, we can implement a translation of our Forge model into native Lean expressions and types.

⁶For example, the top-level definition in Lean “`def x: Int := 0`” is a ‘command’.

⁷Where, + and * denote one/zero or more comma-separated occurrences respectively. + and * denote one/zero or more repetitions.

⁸At least, a useful subset of the Forge language we care about. This is based off the grammar of Alloy [15, 14, 28].

4.3. Elaboration

Elaboration in Lean 4 processes Lean `Syntax` objects, which are the outputs of the Lean parser, into Lean `Expr` objects, which are Lean’s low-level kernel representations [29]. Elaboration is responsible for Lean’s type and metavariable unification⁹, which provides all the type information to Lean.

Expressions also rewrites bound variables using De Bruijn indices.

Complete

4.4. The Forge Model in Lean

Complete, this is mostly a table of Forge concepts and translations in Lean.

⁹That is, types are inferred, coerced, and type classes resolved at this step. Types, including implicit types, and must be fully specified within `Expr` objects.

5. Implementation Details and Challenges

5.1. Finiteness of Forge Sigs

Recall as summarized in [section 4.4](#) that Forge Sigs get (most naturally) translated to Lean Types. However, we need to be cautious using this as a drop-in replacement for the concept of sigs. While we don't have soundness and completeness guarantees, we ought to feel confident that our translation preserves the semantics of Forge faithfully.

One semantic difference in translating Forge Sigs into Lean Types directly is that we lose all sig 'finiteness guarantees' that came with Forge. Since Forge compiles to a bounded SAT problem, all sigs are finitely bounded. This means that we can rely on the assumption that sets of sigs are all finite sets.

For example, we could write a specification of a path through a graph with one source node and claim that this path will never be injective:

```
sig Node {
  next: one Node
}

pred notInjective {
  not ( all a, b: Node |
    a.next = b.next => a = b )
}
```

SK: I'm curious about this statement; it at least deserves elaboration. I see no reason why we can't get some kinds of formal guarantees, since we're going from a formalized object to a formalized object.

Fix this model it's incorrect!

For any number of `Node` sigs we initialize Forge with (that is, for any bandwidth), Forge will be able to tell us that `notInjective` is theorem.

Yet, in the Lean formulation of this specification, we realize that `notInjective` need not be true at all. If we tried to prove the same in Lean, we realize that it isn't actually possible to prove the injectivity of our `next` relation. In fact, if we conjured our type `Node` with the same structure as \mathbb{N} , we would have a perfectly valid `next` function (think `succ`) that is not injective.

The disparity between the Forge and Lean models is that while Forge models are finitely inhabited (albeit, arbitrarily so), Lean makes no assumption about the size of models or types and requires explicit statements of finiteness of types. Pertaining to the example above, we would want to be able to prove the nonexistence of such an injective function using the translated facts produced by our module.

Our solution is to include additional local instance axioms for every opaque sig that is translated from Forge and introduced to our Lean environment. In this case:

```
@[instance] axiom inhabited_node : Inhabited Node
@[instance] axiom fintype_node : Fintype Node
```

which gives us guarantees that `Node` is both inhabited and contains a finite number of elements within the type. To illustrate, our proof of `notInjective` would utilize the pigeon hole principle and appeal to the fact that `Node` is finitely inhabited.

SK: I think there are two things being conflated here: domain bounds and finiteness. Domain bounds obviously imply finiteness but not vice versa.

We discuss further in [section 5.6](#) how these instances enable us to perform integer and cardinality operations on sigs and translated Forge expressions.

5.2. Sig Inheritance and Quantifiers

Many signatures in Forge have complex inheritance structures [15] that cannot be expressed in Lean using a naïve translation. Recall that conventionally, we would translate a signature in Forge into a corresponding opaque type in Lean as follows (see [section 4.4](#)):

```
sig Student {}
```

```
opaque Student : Type
```

However, how then would we represent another sig like `Undergrad` which inherits from a `Student` sig? In Forge, we can use the `extends` keyword to denote that `Undergrad` inherits fields from `Student`:

```
sig Undergrad extends Student {}
```

For `Undergrad` to *extend* `Student`, every field accessible to `Student` must also be available to `Undergrad`, and any expression of the `Undergrad` type should be interchangeable as expressions of type `Student`.

As we did before, we could try to define the corresponding type in the same way, without any regard to the fact that it inherits from such a parent sig:

```
opaque Undergrad : Type
```

However, Lean does not know that all `Undergrad`s are also `Students`, and since fields are typed to the sig that they are a part of in Lean¹⁰, any access into a field that belongs to the `Undergrad` sig inherited because it was part of the `Student` sig would fail. Consider the following Forge program and the wishfully translated Lean equivalent:

```
1 sig Class {}
2 sig Student {
3   registration : set Class
4 }
5 sig Undergrad extends Student {}
6
7 fun ugradsIn[c : Class] : Undergrad {
8   all u : Undergrad |
9     c in u.registration
10 }
```

```
opaque Class : Type
opaque Student : Type
opaque registration : Student → Class → Prop

opaque Undergrad : Type

def ugradsIn (c : Class) : Set Undergrad :=
  ∀ u : Undergrad,
    registration u c -- Type error!
```

Such a Lean translation would raise a type error at line 9 above as `registration` expects an object of the `Student` type for its first input but was given a `Undergrad` type instead. In the case of Forge, Forge is aware that all descendents of a particular sig can be used interchangeably when an expression of that sig is expected. We need to find a (clever) way to encode within Lean that in fact, `Undergrad` is a child sig of `Student` and all `Undergrad`s are `Students`. As Lean is not ‘object-oriented’ in the way that Forge is, there is no directly equivalent concept of a type that inherits from another type in Lean natively.

¹⁰For sig `A` to have a field `f : set A` is for the field `f` to have type `A → A → Prop`.

This task has its subtleties and at the same time we will need to keep usability in mind for an end-user who wishes to prove facts about their model. One direct solution motivated by our type error might be to introduce a coercion instance from `Undergrads` to `Grads` which immediately fixes our problem:

```
@[instance] axiom coe_undergrad_student : Coe Undergrad Student
```

The code snippet above would type check, and we would instantly be able to refer to the child sig in place of its parent sig. However, if we wish to query in a proof whether a `Student` object is an `Undergrad` object as well via a predicate like `IsUndergrad`¹¹, this becomes burdensome and involved:

```
def IsUndergrad (s : Student) : Prop := ∃ x : Undergrad, x = s
```

The existential which quantifies over all `Undergrads` to check if they are equal to `s` is not very user-friendly and can become significantly involved, especially when we are utilizing inheritance liberally in a specific model. Furthermore, we have no straightforward solution given `(IsUndergrad u)` and `(u : Student)` to cast `u` back into an `Undergrad` type.

Instead, we can consider switching the order we define the child type and child type predicate to the dual of the translation above. If instead we define our membership predicate `IsUndergrad` first:

```
opaque IsUndergrad : Student → Prop
```

we can then use Lean’s native subtyping to define our `Undergrad` type:

```
@[reducible] def Undergrad : Type :=
  { s : Student // IsUndergrad s }
```

In this implementation, we also happen to get the `Undergrad` to `Student` coercion automatically as a property of subtyping.

In addition, Forge introduces the concept of *abstract* sigs [15]. If `Student` were an abstract sig, we might encounter `Undergrad` and `Grad` student as concrete subtypes of abstract `Student`. We might encounter:

```
abstract sig Student {}
sig Undergrad extends Student {}
sig Grad extends Student {}
```

which is to say that every object instance of `Student` had ought to be either a `Undergrad` or `Grad`. Within our framework, this can be implemented in Lean as

```
axiom abstract_student : ∀ s : Student, IsUndergrad s ∨ IsGrad s
```

¹¹This is the `Undergrad` membership predicate on `Students`, which we’ll likely need to do in order to prove anything about objects within this inheritance relation.

provided both subtype instances for `Undergrad` and `Grad` have been generated correctly.

Such an analysis and translation of inheritance requires a global processing of the Forge program in Lean, since every expression emitted by our elaborator ought to have well-defined types and terms (that is, cannot be waiting on a term to be defined). Since Forge fields can depend on any sig, we need to ‘lift’ sigs to define first before all fields and predicates can be defined. Furthermore, sig quantifiers and inheritance structures require preprocessing all sigs to generate a topological sort of the inheritance structure before expressing any of their translations in Lean. Since Forge is not a ‘local’ language and all sig declarations are lifted, this was not a problem for us before we entered Lean. As such, sigs can be specified in any order in Forge.

For the remaining of sig quantifiers like `one` and `lone`, because of their complex interactions with inheritance (a `one` sig means that there is only one inhabited member of that sig that *is not* any child sig, contrary to our intuition as to what a `one` or `lone` sig should be), our program prompts the user to write a customized axiom in Lean expressing their desired constraint. Note that this utilizes the seamless integration of Forge into Lean that makes such a solution of mixed execution possible. Anecdotally, `one` and `lone` sigs only apply to child sigs and the `abstract` quantification is only used on parent sigs (it doesn’t make sense for a non-inherited sig to be `abstract`), so we expect that manual handling of sig quantifications to be an edge case.

Write more about types of things here perhaps, separate out into processing for ‘one’ and what not.

5.3. “Everything is a Set”

Overloaded operations

The predominantly relational nature of Forge introduces another point of friction between our translation from Forge to Lean. In Forge, every expression is implicitly a relation or a set (set when that expression has arity-1¹²). Even when we know that an expression is a relation or set with cardinality 1 (for example, it could be introduced as a binder from a quantification), they are used as if they were a singleton set in Forge expressions that expect a set as an operand. Under the hood, all expressions in Forge are treated as a set (or multi-arity relation).

This everything-is-a-set approach of Forge allows the following expression (within the existential quantifier), translated “there is some `Student` who is their own friend”:

```
sig Student {
  friends : set Student
}
pred ownFriend {
  some s: Student |
    s in s.friends
}
```

```
opaque Student : Type
opaque friends : Student → Student → Prop

def ownFriend :=
  ∃ s : Student,
    friends s s
```

Note that `s` is a ‘singleton’, but it is used as if it were an honest-to-goodness set in the join operation (`s.friends`) and the inclusion operation (`s in ...`). We are able to concisely translate into a statement in Lean of the likes of “ $\exists s$ such that on the `friends` relation, $(s, s) \in \text{friends}$.” Note that

¹²The notation we use is that a `Set` α , which has equivalent type $\alpha \rightarrow \text{Prop}$, has arity-1, and so on. This is the arity convention in Forge and aligns with Lean’s definitions of relations.

because s is a singleton—that is, no set with more than one element could possibly be bound to s as a result of our existential quantifier—we were able to translate the join in $s \bowtie \text{friends}$ as the partial application to our relation $\text{friends } s$, and the `in` keyword became set membership $s \in s \bowtie \text{friends}$ which was just $(\text{friends } s) s$.

Consider an alternative when we relax the requirement that s ought to be a singleton in the Forge source:

```
pred ownFriend[t: Student] {
  let s = t.friends |
    s in s.friends
}
```

```
def ownFriend (t : Student) :=
  let s := friends t,
    friends s s -- Type error!
```

Which is loosely “for a `Student t`, the set of t ’s friends is a subset of the set of all *their* friends.” Here, s in Forge is bound to a set of t ’s friends. Had we translated this in the same way, s would be typed $\text{Student} \rightarrow \text{Prop}$ (or equivalently, a `Set Student`), and `friends s s` raise a type error.

We instead have to resolve the join $s \bowtie \text{friends}$ without our shortcut above of partially applying s to `friends`:

```
s.friends
```

```
 $\lambda x_2 \mapsto \exists x_1 : \text{Student}, s\ x_1 \wedge \text{friends } x_1\ x_2$ 
```

which is immediately more cumbersome than our earlier solution.

The same applies when we now try to implement the inclusion `in` operator:

```
s in s.friends
```

```
Set.Subset s ( $\lambda x_2 \mapsto \exists x_1 : \text{Student}, s\ x_1 \wedge \text{friends } x_1\ x_2$ )
```

which becomes a subset operator¹³ instead of set membership.

This example describes a fundamental incompatibility between Forge and Lean that we need to resolve. Forge is indifferent between whether an expression is a singleton or a relation/set, and treats the two indiscriminately. This approach of treating everything as a set allows operations like relational join and ‘in’ to work across all scenarios alike.

However, Lean tends to prefer expressions that are not sets (that is, honest-to-goodness singletons). In the cases above, this allows for relational join to be a partial application, and ‘in’ to be set membership. For the majority of use cases, this singleton-friendly translation suffices. When we are dealing with sets, set operators such as join and ‘in’ (which is now the subset operator) become more convoluted as demonstrated above. Additionally, while joining a singleton and a relation via the partial application solution applies to relations of varying arities, a join expression between two arbitrary relations takes in different types, and hence implementations, depending on the respective arities and types of the arguments.

This means that we shouldn’t take the same approach as Forge of treating everything as a set, and performing the most generic set operation possible on them. Where possible, we ought to keep elements as elements and not cast them into singleton sets, since cutting this corner in translation necessarily comes at the cost that the output of the translation is more complicated.

¹³Under the hood, `Set.Subset s1 s2` is defined as $\forall a, a \in s_1 \rightarrow a \in s_2$.

The outline of our solution is to consistently emit only the simplest (and most type-tailored) translation possible, leveraging the fact that we know at the time of translation all types of the inputs into an operator. We implement this through Lean’s type class system, where a set of methods can be implemented across different types and dispatched according to the type of its arguments. For every pair of types for which a method might be different (in other words, overloaded), we can write an instance of that type class implementing its functionality.

For example, the following is an excerpt¹⁴ of our implementation of relational join as a type class `HJoin` (‘has join’), following our implementations of join from the examples demonstrated above:

```
class HJoin (α : Type) (β : Type) (γ : outParam Type) :=
  (join : α → β → γ)

-- Join singleton with arity-2 relation
@[reducible, simp] instance {α β : Type} : HJoin (α) (α → β → Prop) (β → Prop) where
  join := fun a g ↦ g a

-- Join set with arity-2 relation
@[reducible, simp] instance {α β : Type} : HJoin (α → Prop) (α → β → Prop) (β → Prop) where
  join := fun l r b ↦ ∃ a : α, l a ∧ r a b
```

Then, when translating $a \bowtie b$, we can nondiscriminately produce `HJoin.join a b` and have Lean synthesize which particular implementation to apply based on the types of `a` and `b`. This allows us to have the most specific translation of an expression depending on the types of operands. The `simp` attribute on the instances allows the Lean simplifier (calling the `simp` tactic) to consult this as a unfoldable definition at proof-time and we are left with the native meaning.

For many operators on expressions (see [section 5.5](#) below), their implementations in Lean are overloaded to accomodate the fact that Lean prefers elements when they are elements and sets only when necessary, contrasted to Forge’s ‘everything-is-a-set’ approach. This allows us to produce semantically equivalent translations that are more simplified when possible leveraging Lean’s type class system that is able to determine types of operands at the time of translation.

Types as sets

Forge’s typing ambiguity further exemplifies itself in allowing types to be used where a set is expected to denote the set of all elements in said type. For example:

```
1 pred isAFriend[s: Student] {
2   s in Student.friends
3 }
```

is the predicate that `s` is *someone’s* friend, where the set of all friends is the join of `Student` \bowtie friends (which is equivalent to the set comprehension expression $\{s: \text{Student} \mid \text{some } t: \text{Student} \mid s$

¹⁴There is an instance for every pair of arities and types that could be passed into a join, hence there are many more instances than shown here. However, these implementation details are obscured to the end-user since the join function `HJoin.join` will only resolve to a single instance.

in `t.friends`}). Here, `Student` is being used to denote the *type* of `s` on line 1 and the set corresponding to type `Student` on line 2.

In our translation, we want to be able to use the type `Student` interchangeably in places that expect a set of type `Student` as the set of all `Students`. Since we included our finiteness `Fintype` property as an instance when translating sigs (see [section 5.1](#)), we can create a coercion that coerces a Lean `Type`, given that it is a finite type, into a set of that type using the definition of the `Fintype` typeclass.

check this
after writing
that section

```
instance [f: Fintype α] : CoeDep Type (α : Type) (Set α) where
  coe := (f.elems : Set α)
```

5.4. Type Coercions

complete!

5.5. Relational Joins, Cross, Inclusion, Equality

We need to take special care when implementing expression operators in Lean whenever one of these conditions are true:

- (1) There is no direct out-of-the-box translation for a Forge operation within Lean, or
- (2) there are several implementations of a Forge operation in Lean, depending on the types of the operands given, and where the most generic might not necessarily be the simplest.

We discuss (2) extensively in [section 5.3](#), and introduce using Lean's type class system to implement varying translations of a method depending on the input types. There are several other Forge operations that require this treatment: membership, join (introduced above), cross, and equality.

The specific operators that we needed to take special care translating, the different types that they permit, and their translations in Lean are detailed below in [table 1](#).

Custom syn-
tax for the
binary op-
erations to
pretty print

5.6. Integers

Forge and Alloy come with a unique integer model—due to the fact that models are compiled down to boolean constraints to be solved by a SAT solver, integers are severely limited in their bitwidth [\[15, 27\]](#). The default bitwidth on integers in Forge is 4, which gives us a grand total of 16 integers.

Some programs will inadvertently utilize integer overflow which affects their model in a meaningful ways, oftentimes counterintuitively [\[28, p. 22\]](#). However, the prevailing documentation on integers in the Forge and Alloy solvers casts this effect as a necessary compromise in the design of the language architecture, and placing a bitwidth on integers is an unavoidable consequence due to the boolean formula-based backend of the language.

This is an area where Lean shines. Lean has an integer model which is defined using an honest-to-goodness inductive model for the natural numbers [\[3\]](#). Lean's integer model is arbitrary precision and designed for proofwriting and numerical reasoning. This includes being able to reason about

Forge Operator	Possible Types (singletons lowercase, sets uppercase)	Lean Implementation
Membership: $a \text{ in } b$	$a \text{ in } b$ $a \text{ in } B$ $A \text{ in } b$ $A \text{ in } B$	$a = b$ $a \in B$ $A = \text{Set.singleton } b$ $A \subseteq B$
Equality: $a = b$	$a = b$ $a = B \text{ or } A = b$ $A = B$	$a = b$ $\text{Set.singleton } a = B$, or vice versa. $A = B$
Join: $a.b \text{ or } b[a]$	$a.B$ $A.B$	$B \ a$ <i>Varies, see section 5.3.</i>
Cross: $a \rightarrow b$	$a \rightarrow b$ $a \rightarrow B \text{ or } A \rightarrow b$ $A \rightarrow B$	(a, b) <i>Varies, like $\lambda a f a' b \mapsto a = a' \wedge f b$</i> $\lambda f g a b \mapsto f a \wedge g b$

Table 1.: Forge binary operators and corresponding implementations based on operand types.

integers and write functions involving integers that are noncomputable, such as computing the cardinality of a set.

Here, we depart briefly from the convention that we’ve been following so far of reproducing Forge as accurately as possible in favor of both more extensive integer support, as well as reduced complexity in implementing integers within our translation. For the most part, we are able to use Lean’s integer and finite set/types libraries out of the box with little modification.

Our translated Lean models treats all integers as expressions, which makes the translation from an integer expression in Forge to an integer in Lean relatively straightforward.

Our model is semantically equivalent to running Forge with an arbitrary bitwidth, more than the model would ever exceed and overflow. This gives the most accurate translation of what Forge tries to achieve with integers but is not technically capable of doing.

Here are two examples adapted from [15] that showcase some of the integer features in Forge.

In Forge, the `#` keyword, like on line 4 below, denotes the cardinality of a set.

```

1 sig Suit {}
2 sig Card { suit: one Suit }
3 pred threeOfAKind[hand: set Card] {
4   #hand.suit = 1 and #hand = 3
5 }
```

Fields of sigs can also be integers, and we can do arithmetic on them. By treating all integers as first-class expressions¹⁵, we can also use integers in fields alongside integers that are the result of a set computation. For example, we could define a weighted graph with weighted edges *and* nodes:

```

1  sig Node {
2    node_weight: one Int
3  }
4  sig Edge {
5    start: one Node,
6    end: one Node,
7    edge_weight: one Int
8  }
9  pred nodeWeightIsEdgeWeightPlusOne[n: Node] {
10   n.node_weight = add[1, sum e: { l: Edge | l.start = n } | { e.edge_weight }]
11 }

```

On line 10, we are defining a predicate that states a node n 's weight is equal to 1 plus the sum of edge weights of those edges that start at n .

Of the integer operations, we can easily translate arithmetic operations (addition, subtraction, integer division, remainder, absolute value and sign) as well as inequalities directly into their integer equivalent in Lean. What requires more effort are the notions of counting (cardinality) and quantification in Forge as exemplified above.

We approach the cardinality problem by using `Set.ncard`¹⁶. While this function has a junk value when a set is infinite, we had remedied this earlier in [section 5.1](#) by including `Fintype` axioms with every Forge type we introduce. Lean knows that every set of a `Fintype` is a `Finset` and has an honest-to-goodness cardinality. This allows us to implement `sum`, `max`, `min`, and a summation with a binder (see line 10 of the graph example above) using Lean `Finset` methods such as `Finset.sum`, `Finset.max`, etc.

To illustrate, the translations of the two predicates (playing card hand and graph) above in Lean, eliding sig and field translations, would be as follows:

```

def threeOfAKind (hand : Set Card) : Prop :=
  Set.ncard (hand ⋈ suit) = 1 ∧ Set.ncard hand = 3

```

and

```

def nodeWeightIsEdgeWeightPlusOne (n : Node) : Prop :=
  node_weight n = 1 + Finset.sum { e : Edge | start e = n } edge_weight -- See footnote 17

```

¹⁵Forge denotes integers as atoms or values depending on whether an integer appears in a field or as a result of a computation, but casts seamlessly between [26, 27].

¹⁶More specifically, we do need to utilize the approach in [section 5.3](#) of using type classes to implement this, since the cardinality of a singleton ought to be 1. In all other cases, `Set.ncard` is the implementation of cardinality.

¹⁷There are select details surrounding type coercions, universe levels, and the noncomputability of our integer functions that still remain to be resolved.

While we did need to retrofit additional instance axioms for each type generated to make an integer model work, it is impressive that we were able to extract so much integer functionality out of Forge in within our limited Lean model in the first place. Implementing Forge integers within our translation is also a hallmark of the motivation behind our project in the first place—that in some cases, we are able to endow *additional* functionality to the Forge specification language by interpreting it in a proof assistant instead of the standard relational Forge implementation.

6. Results and Examples

6.1. Forge as a Lean DSL

One of the crucial benefits of working with Lean 4 as a metaprogramming language and a target for our translation is the rich support for DSL implementation and integration. Lean and its accompanying Language Server Protocol (LSP)¹⁸ are designed to have highly flexible and extensible user interfaces that expose useful APIs for implementers of DSLs and custom UI to utilize [24, 25]. Furthermore, Lean’s extensible syntax and macro system are simple yet remarkably powerful [32, 29].

It is as such that we justify framing our implementation of Forge in Lean as an honest-to-goodness domain-specific language (DSL). We do not treat user experience of our tool as an afterthought, nor do we skimp over ensuring that LForge has a set of developer aids just as capable as those found in Forge or Lean themselves. As mentioned in [section 4.1](#), the fact that we can interact with Lean’s implementation means that many of Lean’s ‘IDE-like’ features are exposed to us and available for us to use in the Forge DSL without much overhead.

The following are some (non-exhaustive) examples of the user experience and interface of Forge within Lean.

Syntax Highlighting

Superficially, by defining our syntax as Lean objects and isolating our keywords (we piggyback off Lean’s lexer), we get syntax highlighting of Forge code ‘for free’, on par with Forge’s native solutions. In [fig. 2](#), the syntax of a Forge specification is color-coded.

Types on Hover

Lean exposes an `addTermInfo` method that allows us to attach declared names to pieces of syntax (nodes in Lean’s concrete syntax tree), including custom syntax like ours for Forge. As such, we can annotate relevant pieces of syntax within our Forge specification to reflect names and types that are within scope. When the user hovers their mouse over a piece of syntax corresponding to a Forge expression, a hovering tooltip will display with the type of the expression. In [fig. 2](#), the tooltip shows the type of a Forge predicate defined earlier in the file.

Documentation

As a pedagogical language, Forge has a focus on usability, learnability, and helpful feedback [28], especially when its parent language Alloy is far more permissive with errors. We follow in the same vein in reporting errors and missing features, and in addition, we include documentation on Forge’s syntax within Forge’s on-hover features.

¹⁸This is the language server that processes Lean code and communicates with the code editor or integrated development environment. In this case, we use VS Code.


```

1  import Lforge
2
3  sig Person {
4    | shaved_by: one Person
5  }
6
7  pred shavesThemselves[p: Person] {
8    | p = p.shaved_by
9  }
10
11 pred existsBarber {
12   some b shavesThemselves (p : Person) : Prop
13   not shavesThemselves[p] <=> p.shaved_by = barber
14 }
15 }

```

Figure 2. Tooltips containing type information are available on hover. Forge syntax is automatically highlighted without any extra work.

```

7  pred shavesThemselves[p: Person] {
8    | p = p.shaved_by
9  }
10
11 pred existsBarber {
12   some b shavesThemselves (p : Person) : Prop
13   not shavesThemselves[p] <=> p.shaved_by = barber
14 }
15 }

```

<fmla-a> <=> <fmla-b> : true when <fmla-a> evaluates to true exactly when <fmla-b> evaluates to true. Can also be written as iff . Produces <fmla-a> ↔ <fmla-b> .

```

3  sig Person {
4    shaved_by: one Person
5  }
6
7  Fields
8  Fields allow us to define relationships between a given sig s and other components of our
9  model. Each field in a sig has:
10
11 • a name for the field;
12 • a multiplicity ( one , lone , pfunc , func , or, in Relational or Temporal Forge, set );
13 • a type (a -> separated list of sig names).
14
15 Here is a sig that defines the a Person type with a bestFriend field:
16
17 sig Person {
18   bestFriend: lone Person
19 }
20
21 The lone multiplicity says that the field may contain at most one atom. (Note that this

```

Figure 3. We can define our syntax definitions to print with custom documentation text for users new to using Forge syntax.

Forge documentation is included via docstrings that are placed inline with our syntax objects (see [section 4.2](#)), which is automatically included by Lean’s LSP to display on the front end. [Figure 3](#) showcases docstrings of varying verbosity for operators as well as declaration syntax.

Additionally, we need to be clear and verbose about language features that are not supported in LFORGE. Since LFORGE includes a subset of relational Forge determined by compatibility with Lean’s semantics, we prompt users attempting to use unsupported language features with clarification and a request to redefine their statements. [Figure 4](#) showcases an example of a prompt that lone sig quantifier is unsupported and potentially ambiguous.

Error Checking

Compared to Forge or Racket, Lean (and consequently, LFORGE) provides a markedly better experience with error messages and prompting users when there are errors present in their source program. Since Lean runs in the background as an LSP, users immediately get immediate feedback on whether their source code parses and ‘compiles’.¹⁹

Lean’s error locality system allows its error monad to refer to any piece of syntax object to potentially throw an error. This allows us to prompt errors as soon and as granular as possible. [Figure 5](#) illustrates error reporting at the level of specific identifiers.




Figure 4. We can define custom error messages with our implementation to prompt users to change their specifications if a piece of syntax is ambiguous or not supported.

Types

Lean’s dependent type system is both a blessing and a curse when it comes to the task of translating a language with a foreign type system into Lean. While the type system is a highly optimized algorithm that attempts to resolve type coercions, type classes, and reductions [8], it is often delicate and temperamental, especially when we are working at such a low level of emitting Lean `exprs`, which happens *after* type unification (see [section 4.3](#)). We discuss some of the downsides of such a strict type system in [section 5.4](#), and introduced LFORGE features that circumvent Lean’s restrictions and play into Lean’s type system.

Here, we discuss some of the merits of implementing a DSL designed around Lean’s extensive type system. One of the side effects of translating Forge into Lean is that we inherit Lean’s powerful

¹⁹Forge translations in Lean are not executable, so they provide their value in being interactive with the proof system.

type unification and checking system. This allows us, at specification-time, to check for type errors within the specification. Alloy, on the other hand, is purposefully untyped [14] and only reports type errors at runtime when the successful evaluation of expressions results in the empty expression [10]. This proves difficult to debug and unwieldy for users to understand, as [28] observes. For students who are newly learning the idea of relations, sets, and units, lacking instantaneous feedback on the validity of types and expressions is immensely useful.

Since expressions in our Forge DSL need to translate to typed terms in Lean, we necessarily have to specify types (or use Lean metavariables awaiting unification in place of types) to the Lean expressions emitted. Fortunately, much of this process is abstracted away by the type inference system in place in Lean. For example, to make an application, say, a set union, we don't need to specify the type of set that are being used in the operands and `mkAppM` will complete that for us. However, if we specified two sets of different types, Lean would raise an error. This results in a type-checking and inference system that is as powerful as Lean's with minimal overhead.

Since the Lean LSP provides type linting, our Forge DSL inherits this feature as well. In fig. 5, we pass the `Board` and `Player` arguments to `winRow` in the wrong order, which causes a type error. Lean can identify that the first input to `winRow`, `p`, has the wrong type and displays an appropriate error message.

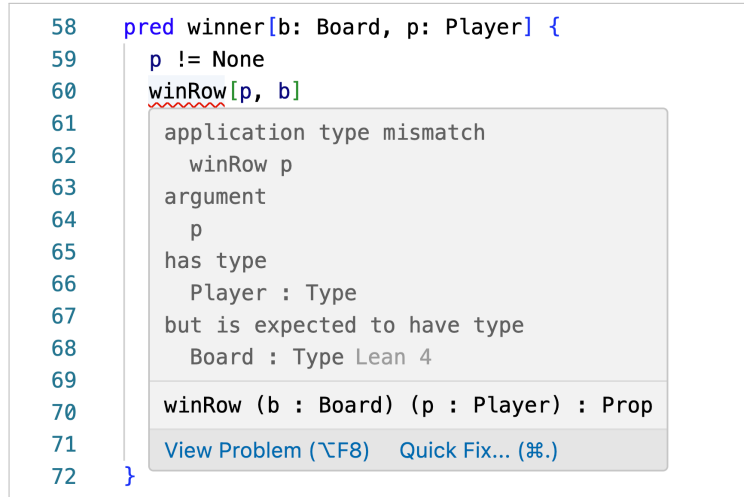


Figure 5. A Lean error message indicating a type mismatch in our Forge expression.

Mixed Execution

Our embedding of Forge within Lean, especially our choice to map Forge structures (sigs, fields, predicates) to corresponding Lean concepts (types, relations, functions) as faithfully as possible means that a Lean file with Forge specifications supports mixed execution, an embedding model explored and supported by similar tools that merge the Alloy specification into other imperative programming languages [22, 23, 21].

All Forge specifications, once inserted into a Lean file, will generate appropriate definitions directly into the Lean environment. Using Lean syntax, we can further interact with these definitions from Forge, perhaps writing predicates or definitions that build on these. This interaction goes the opposite way as well: where Forge expects an expression, predicate, or function, declarations from Lean can be used seamlessly. This provides a frictionless user experience and allows the user to add additional constraints and rules, written in Lean, to a preexisting Forge model. Furthermore, this alleviates the tension incurred of creating the perfect translation: we are aware of the fact that only a subset of Forge is implemented due to the technical restrictions of both platforms. However, with appropriate error reporting (see above Error Checking), users can be prompted to extend their Forge specifications with additional Lean rules. Mixed execution of Forge and Lean means that Forge specifications are now extensible using a much broader functional programming language.

The following is an example of mixed execution of Forge and Lean using LFORGE. The full specification of this example, a model for mutual exclusion of processes, is discussed in [section 6.3](#).

```

1  import Lforge
2
3  abstract sig Location {}
4  one sig Uninterested, Waiting, InCS extends Location {}
5
6  sig Process {}
7
8  sig State {
9    loc: func Process -> Location,
10   flags: set Process
11 }
12
13 def flags_good (s : State) :=
14   ∀ (p : Process), loc s p = InCS → v loc s p = Waiting → flags s p
15
16 pred good[s: State] {
17   flags_good[s]
18   lone {p: Process | s.loc[p] = InCS}
19 }

```

While the majority of this specification is in Forge, we are working in Lean using LFORGE (line 1). We define relevant sigs and fields as a Forge specification. On lines 13-14, we use the types defined in Forge to write a predicate in Lean that states that all processes waiting or in a critical state have a flag raised. In line 17, we've switched back to specifying in Forge but can continue to utilize the `flags_good` predicates we wrote above in Lean. There are no walls or abstractions between the two languages, and Forge is indeed a first-class citizen in the Lean environment.

Since this interoperability works across imports and modules, we envision projects where sections of specifications can be written in Forge and other relevant sections in Lean, allowing users to interoperate between the two. This could also introduce possibilities for users accustomed to the two distinct languages or formalization techniques to collaborate on a shared specification.

6.2. A Toy Example, *Continued*

We revisit our toy example from [section 3.1](#) (the barber who “shaves all those, and only those, who do not shave themselves”). Recall that we had introduced a Forge specification for this problem, as well as an equivalent Lean specification of the paradox. We concluded earlier that while it was insightful for Forge to produce a result that this specification was `Unsatisfiable`, we might still desire a general proof of this fact outside of Forge’s finite and restricted search bounds.

Here’s an example of what the last part of the modeling workflow—writing and completing the proof—would look like in Lean’s interactive tactic mode:

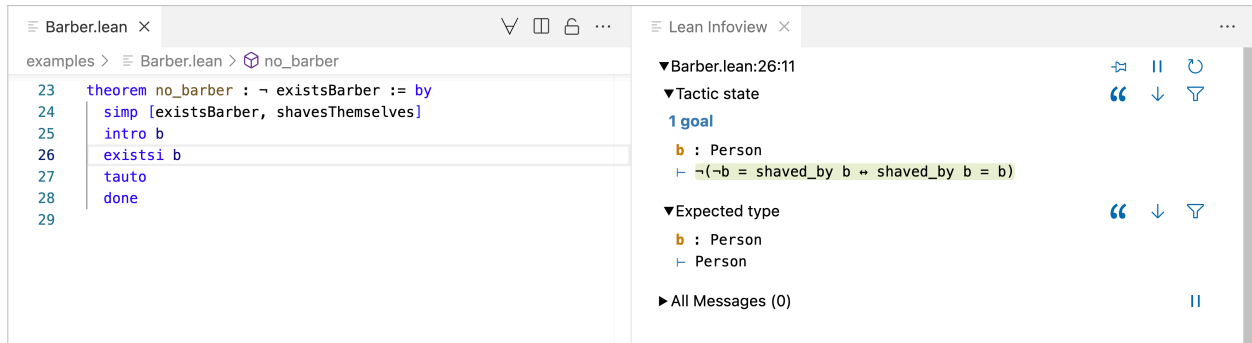


Figure 6. The proof of the nonexistence of a barber in the barber paradox, in Lean. The interactive proof state is on the right with the proof source on the left.

On line 25, we use the `simp` tactic (or we can also use `simp only` or `rw`) to rewrite our Forge-defined predicates `existsBarber` and `shavesThemselves`. Due to the simplicity of the example, the goal can be closed using the `tauto` (tautology) tactic which repeatedly breaks down assumptions and splits goals with with logical connectives until it can close the goal. Full tactic states at each step of this proof are provided in [appendix B](#).

While simple, this example demonstrates the expressiveness of Forge programs embedded in Lean and the relative ease with which some proofs of translated properties can be executed. The following section, [section 6.3](#), presents a more elaborative example of LFORGE.

6.3. A Mutual-Exclusion Protocol

- Circle back to the earlier example
- A more complicated example, potentially with Lean mixins?

7. Discussion

7.1. Contributions

7.2. Reflections

7.3. Future Work

- Summarize contributions.
- Resummarize sections and tie to introduction.
- Future work!

Bibliography

- [1] Jean-Raymond Abrial. “Formal methods in industry: achievements, problems, future”. In: *Proceedings of the 28th international conference on Software engineering*. 2006, pp. 761–768.
- [2] Jeremy Avigad. “Learning logic and proof with an interactive theorem prover”. In: *Proof technology in mathematics research and teaching* (2019), pp. 277–290.
- [3] Jeremy Avigad et al. *Theorem proving in Lean*. 2024.
- [4] Yves Bertot. “A short presentation of Coq”. In: *Theorem Proving in Higher Order Logics: 21st International Conference, TPHOLs 2008, Montreal, Canada, August 18-21, 2008. Proceedings 21*. Springer. 2008, pp. 12–16.
- [5] Avik Chaudhuri et al. “Fast and precise type checking for JavaScript”. In: *Proceedings of the ACM on Programming Languages* 1.OOPSLA (2017), pp. 1–30.
- [6] Albert Mo Kim Cheng. “A survey of formal verification methods and tools for embedded and real-time systems”. In: *International Journal of Embedded Systems* 2.3-4 (2006), pp. 184–195.
- [7] Leonardo De Moura and Nikolaj Bjørner. “Z3: An efficient SMT solver”. In: *International conference on Tools and Algorithms for the Construction and Analysis of Systems*. Springer. 2008, pp. 337–340.
- [8] Leonardo De Moura et al. “The Lean theorem prover (system description)”. In: *Automated Deduction-CADE-25: 25th International Conference on Automated Deduction, Berlin, Germany, August 1-7, 2015, Proceedings 25*. Springer International Publishing. 2015, pp. 378–388.
- [9] Richard St-Denis. “A comparison of three solver-aided programming languages: α Rby, ProB, and Rosette”. In: *Journal of Computer Languages* 77 (2023), p. 101238.
- [10] Jonathan Edwards, Daniel Jackson, and Emina Torlak. “A type system for object models”. In: *ACM SIGSOFT Software Engineering Notes* 29.6 (2004), pp. 189–199.
- [11] George Fink and Matt Bishop. “Property-based testing: a new approach to testing for assurance”. In: *ACM SIGSOFT Software Engineering Notes* 22.4 (1997), pp. 74–80.
- [12] Zheng Gao, Christian Bird, and Earl T Barr. “To type or not to type: quantifying detectable bugs in JavaScript”. In: *2017 IEEE/ACM 39th International Conference on Software Engineering (ICSE)*. IEEE. 2017, pp. 758–769.
- [13] Vladimir Gladshstein, George Pirlea, and Ilya Sergey. “Small Scale Reflection for the Working Lean User”. In: *arXiv preprint arXiv:2403.12733* (2024).

- [14] Daniel Jackson. “Alloy: a language and tool for exploring software designs”. In: *Communications of the ACM* 62.9 (2019), pp. 66–76.
- [15] Daniel Jackson. *Software Abstractions: logic, language, and analysis*. MIT press, 2012.
- [16] Philipp Körner and Florian Mager. “An embedding of B in Clojure”. In: *Proceedings of the 25th International Conference on Model Driven Engineering Languages and Systems: Companion Proceedings*. 2022, pp. 598–606.
- [17] Felix Kossak and Atif Mashkoor. “How to select the suitable formal method for an industrial application: a survey”. In: *International conference on abstract state machines, alloy, b, tla, vdm, and z*. Springer. 2016, pp. 213–228.
- [18] Thierry Lecomte et al. “Applying a formal method in industry: a 25-year trajectory”. In: *Formal Methods: Foundations and Applications: 20th Brazilian Symposium, SBMF 2017, Recife, Brazil, November 29–December 1, 2017, Proceedings 20*. Springer. 2017, pp. 70–87.
- [19] Robert Y Lewis and Minchao Wu. “A bi-directional extensible interface between lean and mathematica”. In: *Journal of Automated Reasoning* 66.2 (2022), pp. 215–238.
- [20] David R MacIver, Zac Hatfield-Dodds, et al. “Hypothesis: A new approach to property-based testing”. In: *Journal of Open Source Software* 4.43 (2019), p. 1891.
- [21] Aleksandar Milicevic et al. “Advancing declarative programming”. PhD thesis. Massachusetts Institute of Technology, 2015.
- [22] Aleksandar Milicevic et al. “Executable specifications for Java programs”. PhD thesis. Massachusetts Institute of Technology, 2010.
- [23] Aleksandar Milicevic, Ido Efrati, and Daniel Jackson. “ α Rby—an embedding of Alloy in Ruby”. In: *Abstract State Machines, Alloy, B, TLA, VDM, and Z: 4th International Conference, ABZ 2014, Toulouse, France, June 2-6, 2014. Proceedings 4*. Springer. 2014, pp. 56–71.
- [24] Leonardo de Moura and Sebastian Ullrich. “The lean 4 theorem prover and programming language”. In: *Automated Deduction—CADE 28: 28th International Conference on Automated Deduction, Virtual Event, July 12–15, 2021, Proceedings 28*. Springer. 2021, pp. 625–635.
- [25] Wojciech Nawrocki, Edward W Ayers, and Gabriel Ebner. “An extensible user interface for Lean 4”. In: *14th International Conference on Interactive Theorem Proving (ITP 2023)*. Schloss Dagstuhl-Leibniz-Zentrum für Informatik. 2023.
- [26] Tim Nelson. *Forge Language Documentation*. <https://cscil710.github.io/forge-documentation/home.html>. 2024.
- [27] Tim Nelson et al. *Artifact for Forge: A Tool and Language for Teaching Formal Methods*. Version 0.3. Jan. 2024. DOI: [10.5281/zenodo.10463960](https://doi.org/10.5281/zenodo.10463960). URL: <https://doi.org/10.5281/zenodo.10463960>.
- [28] Tim Nelson et al. “Forge: A Tool and Language for Teaching Formal Methods”. In: *Proceedings of the ACM on Programming Languages* 8.OOPSLA1 (2024), pp. 1–31.

- [29] Arthur Paulino et al. *Metaprogramming in Lean 4*. 2024. URL: <https://leanprover-community.github.io/lean4-metaprogramming-book/>.
- [30] Bertrand Russell. *The philosophy of logical atomism*. Routledge, 2009.
- [31] Emina Torlak and Daniel Jackson. “Kodkod: A relational model finder”. In: *International Conference on Tools and Algorithms for the Construction and Analysis of Systems*. Springer. 2007, pp. 632–647.
- [32] Sebastian Ullrich and Leonardo De Moura. “Beyond notations: Hygienic macro expansion for theorem proving languages”. In: *Logical Methods in Computer Science* 18 (2022).
- [33] Jim Woodcock et al. “Formal methods: Practice and experience”. In: *ACM computing surveys (CSUR)* 41.4 (2009), pp. 1–36.

Appendices

A. Data Availability

The source code for LFORGE, including all mentioned examples and proofs, is available at this repository: <https://github.com/jchen/lforge>. The source code at the time of this thesis being submitted is tagged `thesis`. LFORGE is available as a package and can be included as a dependency using Lake.

[Check tag.](#)

B. Barber Paradox Proof

The proof of the barber paradox annotated with the Lean tactic state after each tactic/step is provided below. This is also provided in a file at this path in the code repository: `examples/Barber.lean`.

[check
filepath](#)

```
theorem no_barber : ¬ existsBarber := by
  /-
  ⊢ ¬existsBarber
  -/
  simp [existsBarber, shavesThemselves]
  /-
  ⊢ ∀ (x : Person), ∃ x_1, ¬(¬x_1 = shaved_by x_1 ↔ shaved_by x_1 = x)
  -/
  intro b
  /-
  b : Person
  ⊢ ∃ x, ¬(¬x = shaved_by x ↔ shaved_by x = b)
  -/
  existsi b
  /-
  b : Person
  ⊢ ¬(¬b = shaved_by b ↔ shaved_by b = b)
  -/
  tauto
done
```